

NeuralAssignment-7.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

RAM Disk

+ Code + Text

Simple CNN model for CIFAR-10

import numpy

from keras.datasets import cifar10

from keras.models import Sequential

from keras.layers import Dense

from keras.layers import Dropout

from keras.layers import Flatten

from keras.constraints import MaxNorm

from keras.optimizers import SGD

from keras.layers import Conv2D

from keras.layers import MaxPooling2D

from keras.utils import to_categorical

from keras import backend as K

set image dim ordering to 'th'

fix random seed for reproducibility

seed = 7

numpy.random.seed(seed)

load data

(X_train, y_train), (X_test, y_test) = cifar10.load_data()

normalize inputs from 0-255 to 0.0-1.0

X_train = X_train.astype('float32')

X_test = X_test.astype('float32')

X_train = X_train / 255.0

X_test = X_test / 255.0

one hot encode outputs

y_train = to_categorical(y_train)

y_test = to_categorical(y_test)

num_classes = y_test.shape[1]

Create the model

model = Sequential()

model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu'))

model.add(Dropout(0.2))

model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))

model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())

model.add(Dense(1024, activation='relu'))

model.add(Dropout(0.2))

model.add(Dense(num_classes, activation='softmax'))

Compile model

epochs = 5

lr_rate = 0.01

decay = lr_rate/epochs

1s completed at 7:38 PM

NeuralAssignment-7.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

Download

RAM Disk

1s

completed at 7:38 PM

```
decay = 1/rste/epochs
sgd = SGD(lr=lr*decay)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())
# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)
# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))
```

Downloading data from <https://www.cs.toronto.edu/~cifar/cifar-10-python.tar.gz>
170408071/170408071 [*****] - 2s 0us/step
WARNING: absl: 'lr' is deprecated in Keras optimizer, please use 'learning_rate' or use the legacy optimizer, e.g., tf.keras.optimizers.legacy.SGD.
Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 128)	36992
max_pooling2d (MaxPooling2D)	(None, 16, 16, 128)	0
flatten (Flatten)	(None, 32768)	0
dense (Dense)	(None, 1024)	33555456
dropout_1 (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 10)	10250

=====
Total params: 33603584 (128.19 MB)
Trainable params: 33603584 (128.19 MB)
Non-trainable params: 0 (0.00 Byte)

```
None
Epoch 1/5
1563/1563 [*****] - 615s 303ms/step - loss: 1.8202 - accuracy: 0.3492 - val_loss: 1.6758 - val_accuracy: 0.4187
Epoch 2/5
1563/1563 [*****] - 611s 391ms/step - loss: 1.5449 - accuracy: 0.4543 - val_loss: 1.4517 - val_accuracy: 0.4877
Epoch 3/5
1563/1563 [*****] - 614s 393ms/step - loss: 1.3696 - accuracy: 0.5135 - val_loss: 1.2700 - val_accuracy: 0.5488
```

NeuralAssignment-7.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share

RAM Disk

+ Code + Text

1563/1563 [-----] - 618s 395ms/step - loss: 1.2337 - accuracy: 0.5643 - val_loss: 1.1732 - val_accuracy: 0.5884
Epoch 5/5
1563/1563 [-----] - 614s 393ms/step - loss: 1.2342 - accuracy: 0.6000 - val_loss: 1.1265 - val_accuracy: 0.6034
Accuracy: 60.34%

[2] import numpy as np
predictions = model.predict(X_test[:4])
predicted_labels = np.argmax(predictions, axis=1)
actual_labels = np.argmax(y_test[:4], axis=1)

print("Predicted labels:", predicted_labels)
print("Actual labels: ", actual_labels)


1/1 [-----] - 0s 141ms/step
Predicted labels: [3 1 8 0]
Actual labels: [3 8 8 0]

[3] import matplotlib.pyplot as plt

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='upper right')
plt.show()

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['train', 'val'], loc='lower right')
plt.show()

Model Loss



Epoch	train	val
0	1.82	1.75
1	1.78	1.72
2	1.75	1.68
3	1.72	1.65
4	1.70	1.63
5	1.68	1.62

1s completed at 7:38 PM

